

## MILESTONE 4: Account Provisioning

· Add each team member's user account to the proper groups to grant them the necessary access. I added users to the "slurm" group, allowing them to submit jobs to Slurm.

```
sudo usermod -aG slurm venkata
sudo usermod -aG slurm kush
sudo usermod -aG slurm roja
sudo usermod -aG slurm hemanth
```

```
saiteja — venkata@t-green: /a/apps/tb/slurm-green-1$ sudo usermod -aG slurm venkata
venkata@t-green:/a/apps/tb/slurm-green-1$ sudo usermod -aG slurm kush
venkata@t-green:/a/apps/tb/slurm-green-1$ sudo usermod -aG slurm hemanth
venkata@t-green:/a/apps/tb/slurm-green-1$ sudo usermod -aG slurm kush01
venkata@t-green:/a/apps/tb/slurm-green-1$ sudo usermod -aG slurm roja
venkata@t-green:/a/apps/tb/slurm-green-1$ cat /etc/group | grep "slurm"
sudo:x:27:ubuntu,yhuang,klam,hlu,hniu,cxu,zgong,sadmin,slurm,venkata
slurm:x:8051:venkata,kush,hemanth,kush01,roja
```

· This command adds a new user account named "c74" to the SLURM accounting system, with a description of "class", associated with the organization "neu", and allowed to use the cluster "green".

```
sacctmgr -i add account c74 Description=class Organization=neu cluster=green
```

```
slurm@t-green:/home/venkata$ sacctmgr -i add account c74 Description=class Organization=neu cluster=green
Associations
A = c74          C = green
Settings
Parent          = root
slurm@t-green:/home/venkata$
```

Now we see the accounts and clusters using the following commands below:

```
sacctmgr show account
sacctmgr show clusters
```

```
venkata@t-green:~$ sacctmgr show account
Account      Descr      Org
-----
c74          class     coe
c74x23      2023 spring class   neu
nandikonda  2023 spring class   neu
none        class svc accounts  none
root default root account  root
test        class svc accounts  none
venkata@t-green:~$ sacctmgr show clusters
Cluster      ControlHost ControlPort  RPC  Share  GrpJobs  GrpTRES  GrpSubmit  MaxJobs  MaxTRES  MaxSubmit  MaxWall  QOS  Def QOS
-----
black        172.31.86.98  6817  9728  1      1      cpu=10      1      1      1      1      normal
blue         172.31.86.96  6817  9728  1      1      1      1      1      1      1      normal
brown        172.31.86.97  6817  9728  1      1      1      1      1      1      1      normal
green        172.31.86.97  6817  9728  1      1      1      1      1      1      1      normal
p01          172.31.89.108 6817  9728  1      1      1      1      1      1      1      normal
venkata@t-green:~$
```

```
class=(kush,kush01,roja,venkata,hemanth)
for u in "${class[@]}"; do
    #echo "user=$u"
    #sacctmgr -i delete account $u
    sacctmgr -i delete user $u
    #echo sacctmgr -i delete user $u
    sacctmgr -i add user $u Accounts+=c74 DefaultAccount=c74 cluster=green defaultqos=normal qos=normal,green
    #sacctmgr -i update user $u set AdminLevel=Admin
    #echo sacctmgr -i update user $u set defaultqos=normal qos=normal
    #echo sacctmgr -i update user $u set defaultqos=p1q1 cluster=p01
done
```

The above script creates or modifies SLURM user accounts for users named "kush", "kush01", "roja", "venkata", and "hemanth". It first deletes any existing accounts for those users, then creates new accounts with the specified account name, default account, cluster, and quality of service settings. It also includes some commented-out lines that can be used to update or modify account settings for these users.

```
slurm@t-green:/home/venkata$ for u in "${class[@]}"; do
> #echo "user=$u"
> #sacctmgr -i delete account $u
> sacctmgr -i delete user $u
> #echo sacctmgr -i delete user $u
> sacctmgr -i add user $u Accounts+=c74 DefaultAccount=c74 cluster=green defaultqos=normal qos=normal,green
> #sacctmgr -i update user $u set AdminLevel=Admin
> #echo sacctmgr -i update user $u set defaultqos=normal qos=normal
> #echo sacctmgr -i update user $u set defaultqos=plq1 cluster=p01
> done
Nothing deleted
Adding User(s)
  kush
  kush01
  roja
  venkata
  hemanth
Settings =
  Default Account = c74
Associations =
  U = kush      A = c74      C = green
  U = kush01    A = c74      C = green
  U = roja      A = c74      C = green
  U = venkata   A = c74      C = green
  U = hemanth   A = c74      C = green
Non Default Settings
  QOS           = +green,+normal
  DefQOS        = normal
slurm@t-green:/home/venkata$
```

We see the account users using the command below:

```
slurm@t-green:/home/venkata$ sacctmgr show user
```

User	Def Acct	Admin
dhruv		Administ+
hemanth	c74	Administ+
hniu		None
kush	c74	Administ+
kush01	c74	Administ+
manoj		None
pawan		Administ+
prem		Administ+
roja	c74	Administ+
root	root	Administ+
varun		Administ+
venkata	c74	Administ+
yhuang		Administ+

```
slurm@t-green:/home/venkata$
```

## MILESTONE 5

Integrate the Slurm controller and worker nodes

`sinfo -N -o "%.20N %.15C %.8t %.12m %.12P %10f %.15G %.40E"`

```
slurm@t-green:/home/venkata$ sinfo -N -o "%.20N %.15C %.8t %.12m %.12P %10f %.15G %.40E"
NODELIST      CPUS(A/I/O/T)  STATE  MEMORY  PARTITION  AVAIL_FEAT  GRES  REASON
i-hemanth     0/1/0/1        idle    1        batch* (null)  (null)  (null)  none
i-hemanth     0/1/0/1        idle    1        ig0 (null)    (null)  (null)  none
i-kush         0/1/0/1        idle    1        batch* (null)  (null)  (null)  none
i-kush         0/1/0/1        idle    1        ig0 (null)    (null)  (null)  none
i-roja         0/1/0/1        unk*    1        ig0 (null)    (null)  (null)  none
i-roja         0/1/0/1        unk*    1        batch* (null)  (null)  (null)  none
i-venkata     0/1/0/1        unk*    1        ig0 (null)    (null)  (null)  none
i-venkata     0/1/0/1        unk*    1        batch* (null)  (null)  (null)  none
slurm@t-green:/home/venkata$
```

Test the integrated system and ensure it is functioning properly

Testing from team servers to all individual nodes

`srun -p ig0 --nodes=4 hostname`

```
venkata@t-green:~$ srun -p ig0 --nodes=4 hostname
i-kush
i-venkata
i-roja
i-hemanth
venkata@t-green:~$
```

Individual Ping screenshots to team server

1.

```
venkata@i-venkata:~$ scontrol ping
Slurmctld(primary) at t-green is UP
venkata@i-venkata:~$
```

2.

```
kush01@i-kush:~$ scontrol ping
Slurmctld(primary) at t-green is UP
kush01@i-kush:~$
```

```
hemanth@i-hemanth:~$ scontrol ping
Slurmctld(primary) at t-green is UP
hemanth@i-hemanth:~$
```

3.

```
roja@i-roja:~$ scontrol ping
Slurmctld(primary) at t-green is UP
roja@i-roja:~$
```

4.

## MILESTONE 6: Optimization, scaling and Qos

This command adds a new Quality of Service (QoS) named "green" to the SLURM workload manager. The QoS specifies a maximum of 1 CPU per user, a priority of 10, and a "DenyOnLimit" flag that prevents jobs from running once the specified limit is reached. Additionally, this QoS has a maximum walltime limit of 5 minutes.

sacctmgr -i add qos green set MaxTRESPerUser=cpu=1 priority=10 Flags=DenyOnLimit MaxWall=0-0:5:0

```
slurm@t-green:/home/venkata$ sacctmgr -i add qos green set MaxTRESPerUser=cpu=1 priority=10 Flags=DenyOnLimit MaxWall=0-0:5:0
Adding QoS(s)
green
Settings
Description    = green
Flags          = DenyOnLimit
MaxTRESPerUser = cpu=1
MaxWall        = 00:05:00
Priority       = 10
```

This command displays a list of all Quality of Service (QoS) configurations that have been defined in the SLURM workload manager. It provides information such as the QoS name, description, and resource allocation settings, including limits on CPU usage, memory usage, and run time.

sacctmgr show qos

```
venkata@t-green:~$ sacctmgr show qos
Name      Priority  GraceTime  Preempt  PreemptExemptTime  PreemptMode  MaxTRESPerUser  MaxJobsPU  MaxSubmitPU  Flags  UsageThres  UsageFactor  GrpTRES  GrpTRESMins  GrpTRESRunMin  GrpJobs  GrpSubmit
t  GrpWall  MaxTRES  MaxTRESPerNode  MaxTRESMins  MaxWall  MaxTRESPU  MaxJobsPU  MaxSubmitPU  MaxTRESPA  MaxJobsPA  MaxSubmitPA  MinTRES  GrpTRESMins  GrpTRESRunMin  GrpJobs  GrpSubmit
-----
normal    0  00:00:00  cluster  00:05:00  cluster  cpu=1  1  1  1  1.000000  1.000000  1.000000
piql      10  00:00:00  cluster  00:05:00  cluster  cpu=1  1  1  1  DenyOnLimit  1.000000  1.000000  1.000000
green     10  00:00:00  cluster  00:05:00  cluster  cpu=1  1  1  1  DenyOnLimit  1.000000  1.000000  1.000000
black     10  00:00:00  cluster  00:05:00  cluster  cpu=1  1  1  1  DenyOnLimit  1.000000  1.000000  1.000000
venkata@t-green:~$
```

## QOS Limitations

**squeue** is a command-line utility in the SLURM workload manager that displays information about jobs running or waiting to run in a SLURM-managed cluster. It provides a real-time view of job status, including job IDs, user IDs, partition names, start and end times, and resource allocations. The output of the squeue command can be sorted and filtered in various ways to display only the desired information, such as jobs from a specific user, or jobs running on a particular node or partition. The squeue command is a useful tool for monitoring and managing jobs in a SLURM-managed HPC environment.

```
venkata@t-green:/home/venkata/test/slurm_arr$ sbatch sub_arr.sh
sbatch: error: QOSMaxCpuPerUserLimit
sbatch: error: Batch job submission failed: Job violates accounting/QoS policy (job submit limit, user's size and/or time limits)
venkata@t-green:/home/venkata/test/slurm_arr$ squeue
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
venkata@t-green:/home/venkata/test/slurm_arr$ cat sub_arr.sh
#!/bin/bash
#SBATCH --job-name=arrayTest
#SBATCH --nodes=2
#SBATCH --ntasks=4
#SBATCH --time=5:00
#SBATCH --mem=2mb # Job memory request on each node
#SBATCH --output=j_%A_%a.out
#SBATCH --error=j_%A_%a.err
#SBATCH --partition batch
#SBATCH --array=1-4

y=$((SLURM_ARRAY_TASK_ID+2));
echo "$SLURM_ARRAY_TASK_ID" $y
sleep 10

# sbatch sub_arr.sh
```

The first command updates the configuration of the "green" Quality of Service (QoS) in the SLURM workload manager to change the maximum number of CPUs that can be used per user from 1 to 4.

The second command displays a list of all QoS configurations in the SLURM workload manager, including the updated "green" QoS with the new CPU limit.

In summary, these commands modify the resource allocation settings of the "green" QoS and then display the updated QoS configuration.

```
sacctmgr -i update qos green set MaxTRESPerUser=cpu=4
sacctmgr show qos
```

```
venkata@t-green:/home/venkata/test/slurm_arr$ sacctmgr show qos
t Name Priority GraceTime Preempt PreemptExemptTime PreemptMode
GrpWall MaxTRES MaxTRESPerNode MaxTRESMins MaxWall MaxTRESPU MaxJobsPU MaxSubmitPU
Flags UsageThres UsageFactor
MaxTRESPA MaxJobsPA MaxSubmitPA
GrpTRES GrpTRESMins GrpTRESRunMin GrpJobs GrpSubmi
-----
normal 0 00:00:00 cluster 1.000000
piq1 10 00:00:00 cluster 00:05:00 cpu=1 DenyOnLimit 1.000000
green 10 00:00:00 cluster 00:05:00 cpu=4 DenyOnLimit 1.000000
black 10 00:00:00 cluster 00:05:00 cpu=1 DenyOnLimit 1.000000
```

## MILESTONE 7: Advanced Parallel Jobs

### Task 1: Job array

We run these commands for the sub\_arr.sh files

```
[ -d ~/test/slurm_arr ] | mkdir -p ~/test/slurm_arr
cd ~/test/slurm_arr
cp /a/util/demo/slurm/run/slurm_arr/sub_arr.sh ~/test/slurm_arr/
```

```
#!/bin/bash
#SBATCH --job-name=arrayTest
#SBATCH --nodes=4
#SBATCH --ntasks=4
#SBATCH --time=5:00
#SBATCH --mem=2mb # Job memory request on each node
#SBATCH --output=j_%A_%a.out
#SBATCH --error=j_%A_%a.err
#SBATCH --partition batch
#SBATCH --array=1-4

y=$(( $SLURM_ARRAY_TASK_ID+2 ));
echo "$SLURM_ARRAY_TASK_ID" $y
sleep 10

# sbatch sub_arr.sh
```

This command submits a job to the SLURM workload manager using the shell script sub\_arr.sh. The sbatch command reads the script and submits it to the SLURM controller for scheduling and execution. The script sub\_arr.sh will contain instructions for the SLURM controller to allocate resources for the job, such as the number of CPU cores, memory requirements, and expected run time.

Once the job is submitted, the sbatch command will return a job ID number that can be used to track the status of the job using other SLURM commands, such as squeue or sacct.

sbatch sub\_arr.sh

```
venkata@t-green:/home/venkata/test/slurm_arr$ [ -d ~/test/slurm_arr ] | mkdir -p ~/test/slurm_arr
venkata@t-green:/home/venkata/test/slurm_arr$ cd ~/test/slurm_arr
venkata@t-green:/home/venkata/test/slurm_arr$ cp /a/util/demo/slurm/run/slurm_arr/sub_arr.sh ~/test/slurm_arr/
venkata@t-green:/home/venkata/test/slurm_arr$ vi sub_arr.sh
venkata@t-green:/home/venkata/test/slurm_arr$ sbatch sub_arr.sh
Submitted batch job 295
venkata@t-green:/home/venkata/test/slurm_arr$ squeue
      JOBID PARTITION    NAME   USER  ST       TIME  NODES NODELIST(REASON)
      295_[2-4]    batch arrayTes venkata PD          0:00      4 (Resources)
      295_1      batch arrayTes venkata R          0:03      4 i-hemanth,i-kush,i-roja,i-venkata
venkata@t-green:/home/venkata/test/slurm_arr$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0        up    infinite     4    alloc i-hemanth,i-kush,i-roja,i-venkata
batch*     up    infinite     4    alloc i-hemanth,i-kush,i-roja,i-venkata
venkata@t-green:/home/venkata/test/slurm_arr$ squeue
      JOBID PARTITION    NAME   USER  ST       TIME  NODES NODELIST(REASON)
      295_4      batch arrayTes venkata PD          0:00      4 (Resources)
      295_3      batch arrayTes venkata R          0:07      4 i-hemanth,i-kush,i-roja,i-venkata
venkata@t-green:/home/venkata/test/slurm_arr$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0        up    infinite     4    alloc i-hemanth,i-kush,i-roja,i-venkata
batch*     up    infinite     4    alloc i-hemanth,i-kush,i-roja,i-venkata
venkata@t-green:/home/venkata/test/slurm_arr$ squeue
      JOBID PARTITION    NAME   USER  ST       TIME  NODES NODELIST(REASON)
```

sacct is a command-line tool used to display accounting information for SLURM jobs. The `sacct -j` option followed by a job ID number can be used to display detailed information about a specific job.

```
sacct -j 295
```

```
--format=user,jobid,jobname,Partition,NodeList,ntasks,reqtres%30,submit,start,end,elapsed,TimeLimit,ReqMem,State
```

User	JobID	JobName	Partition	NodeList	NTasks	ReqTRES	Submit	Start	End	Elapsed	TimeLimit	ReqMem	State
venkata	295_1	arrayTest	batch	i-hemanth,i-ku-	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:10:41	2023-04-20T20:10:41	2023-04-20T20:10:51	00:00:10	00:05:00	8M	COMPLETED
venkata	295_1.batch	batch	batch	i-hemanth	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:10:41	2023-04-20T20:10:41	2023-04-20T20:10:51	00:00:10	00:05:00	8M	COMPLETED
venkata	295_2	arrayTest	batch	i-hemanth,i-ku-	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:10:52	2023-04-20T20:10:52	2023-04-20T20:11:02	00:00:10	00:05:00	8M	COMPLETED
venkata	295_2.batch	batch	batch	i-hemanth	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:10:52	2023-04-20T20:10:52	2023-04-20T20:11:02	00:00:10	00:05:00	8M	COMPLETED
venkata	295_3	arrayTest	batch	i-hemanth,i-ku-	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:11:03	2023-04-20T20:11:03	2023-04-20T20:11:13	00:00:10	00:05:00	8M	COMPLETED
venkata	295_3.batch	batch	batch	i-hemanth	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:11:03	2023-04-20T20:11:03	2023-04-20T20:11:13	00:00:10	00:05:00	8M	COMPLETED
venkata	295_4	arrayTest	batch	i-hemanth,i-ku-	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:11:14	2023-04-20T20:11:14	2023-04-20T20:11:25	00:00:11	00:05:00	8M	COMPLETED
venkata	295_4.batch	batch	batch	i-hemanth	1	billing=4,cpu=4,mem=8M,node=4	2023-04-20T20:11:14	2023-04-20T20:11:14	2023-04-20T20:11:25	00:00:11	00:05:00	8M	COMPLETED

```
ls -al j_295*
```

```
venkata@t-green:/home/venkata/test/slurm_arr$ ls -al j_295*
-rw-rw-r-- 1 venkata venkata 0 Apr 20 20:10 j_295_1.err
-rw-rw-r-- 1 venkata venkata 4 Apr 20 20:10 j_295_1.out
-rw-rw-r-- 1 venkata venkata 0 Apr 20 20:10 j_295_2.err
-rw-rw-r-- 1 venkata venkata 4 Apr 20 20:10 j_295_2.out
-rw-rw-r-- 1 venkata venkata 0 Apr 20 20:11 j_295_3.err
-rw-rw-r-- 1 venkata venkata 4 Apr 20 20:11 j_295_3.out
-rw-rw-r-- 1 venkata venkata 0 Apr 20 20:11 j_295_4.err
-rw-rw-r-- 1 venkata venkata 4 Apr 20 20:11 j_295_4.out
venkata@t-green:/home/venkata/test/slurm_arr$
```

```
cat j_295_*.out
```

```
venkata@t-green:/home/venkata/test/slurm_arr$ cat j_295_*.err
venkata@t-green:/home/venkata/test/slurm_arr$ cat j_295_*.out
1 3
2 4
3 5
4 6
venkata@t-green:/home/venkata/test/slurm_arr$
```

## Task 2: Chain submission

```
vi chain.sh
```

```
F=$(sbatch --parsable sub1.sh)&& sbatch --dependency=afterok:${F} sub2.sh
```



vi sub1.sh

```
#!/bin/bash
#SBATCH --job-name=first          # Job name
#SBATCH --partition=batch         # send job to c0 partition
#SBATCH --ntasks=1               # Ask for 1 CPU
#SBATCH --time=00:02:00          # Time limit hrs:min:sec
#SBATCH --mem=2mb                 # Job memory request on each node
#SBATCH --output=j_f_%j.out      # Standard output and error log
#SBATCH --error=j_f_%j.err       # Standard output and error log

cd ~/test/slurm_chain

ml Python/3.9.6-GCCcore-11.2.0
rm first.txt

sleep 1m

python -c "with open('first.txt', 'w') as fw: s=f\"first job, job ID is {${SLURM_JOB_ID}}\\n\\n\"; fw.write( s); "

# sbatch sub1.sh
# squeue
```

vi sub2.sh

```
#!/bin/bash
#SBATCH --job-name=second        # Job name
#SBATCH --partition=batch        # send job to c0 partition
#SBATCH --ntasks=1              # Ask for 1 CPU
#SBATCH --time=00:05:00         # Time limit hrs:min:sec
#SBATCH --mem=2mb               # Job memory request on each node
#SBATCH --output=j_s_%j.out     # Standard output and error log
#SBATCH --error=j_s_%j.err      # Standard output and error log

cd ~/test/slurm_chain

ml Python/3.9.6-GCCcore-11.2.0
python -c "fo=open('first.txt', 'r+');line=fo.read();print (\"Read Line: %s\" % (line));fo.close()"
```

```
[ -d ~/test/slurm_chain ] || mkdir -p ~/test/slurm_chain
cd ~/test/slurm_chain
cp -r /a/util/demo/slurm/run/slurm_chain/ ~/test/slurm_chain/
bash chain.sh
```



```

venkata@t-green:/home/venkata/test/slurm_arr$ [ -d ~/test/slurm_chain ] | mkdir -p ~/test/slurm_chain
venkata@t-green:/home/venkata/test/slurm_arr$ cd ~/test/slurm_chain
venkata@t-green:/home/venkata/test/slurm_chain$ cp /a/util/demo/slurm/run/slurm_chain/chain.sh ~/test/slurm_chain/
venkata@t-green:/home/venkata/test/slurm_chain$ cp -r /a/util/demo/slurm/run/slurm_chain/ ~/test/slurm_chain/
venkata@t-green:/home/venkata/test/slurm_chain$ ls -al
total 52
drwxr-xr-x 3 venkata venkata 6144 Apr 20 20:15 .
drwxrwxr-x 7 venkata venkata 6144 Apr 20 14:40 ..
-rw-r--r-- 1 venkata venkata 74 Apr 20 20:15 chain.sh
-rw-rw-r-- 1 venkata venkata 25 Apr 20 14:49 first.txt
-rw-rw-r-- 1 venkata venkata 0 Apr 20 14:48 j_f_280.err
-rw-rw-r-- 1 venkata venkata 0 Apr 20 14:48 j_f_280.out
-rw-r--r-- 1 venkata venkata 57 Apr 20 14:40 j_f_37.err
-rw-r--r-- 1 venkata venkata 0 Apr 20 14:40 j_f_37.out
-rw-r--r-- 1 venkata venkata 0 Apr 20 14:40 j_s_38.err
-rw-r--r-- 1 venkata venkata 36 Apr 20 14:40 j_s_38.out
drwxr-xr-x 2 venkata venkata 6144 Apr 20 20:15 slurm_chain
-rw-r--r-- 1 venkata venkata 632 Apr 20 14:45 sub1.sh
-rw-r--r-- 1 venkata venkata 574 Apr 20 14:45 sub2.sh
venkata@t-green:/home/venkata/test/slurm_chain$ vi chain.sh
venkata@t-green:/home/venkata/test/slurm_chain$ vi sub1.sh
venkata@t-green:/home/venkata/test/slurm_chain$ vi sub2.sh
venkata@t-green:/home/venkata/test/slurm_chain$ bash chain.sh
Submitted batch job 300
venkata@t-green:/home/venkata/test/slurm_chain$ squeue
      JOBID PARTITION    NAME    USER  ST       TIME  NODES NODELIST(REASON)
      300      batch    second  venkata  PD       0:00      1 (Dependency)
      299      batch    first   venkata  R       0:03      1 i-hemanth
venkata@t-green:/home/venkata/test/slurm_chain$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0        up    infinite     1   alloc i-hemanth
ig0        up    infinite     3   idle  i-kush,i-roja,i-venkata
batch*     up    infinite     1   alloc i-hemanth
batch*     up    infinite     3   idle  i-kush,i-roja,i-venkata
venkata@t-green:/home/venkata/test/slurm_chain$ █

```

sacct -j 299,300

--format=user,jobid,jobname,Partition,NodeList,ntasks,reqtres%30,submit,start,end,elapsed,TimeLimit,ReqMem,State

```

venkata@t-green:/home/venkata/test/slurm_chain$ sacct -j 299,300 --format=user,jobid,jobname,Partition,NodeList,ntasks,reqtres%30,submit,start,end,elapsed,TimeLimit,ReqMem,State
User JobID JobName Partition NodeList NTasks ReqTRES Submit Start End Elapsed TimeLimit ReqMem State
venkata 299 first batch i-hemanth 1 billing=1,cpu=1,mem=2M,node=1 2023-04-20T20:16:00 2023-04-20T20:16:00 2023-04-20T20:17:03 00:01:03 00:02:00 2M COMPLETED
venkata 299.batch batch batch i-hemanth 1 2023-04-20T20:16:00 2023-04-20T20:16:00 2023-04-20T20:17:03 00:01:03 00:02:00 2M COMPLETED
venkata 300 second batch i-hemanth 1 billing=1,cpu=1,mem=2M,node=1 2023-04-20T20:16:00 2023-04-20T20:17:03 2023-04-20T20:17:04 00:00:01 00:05:00 2M COMPLETED
venkata 300.batch batch batch i-hemanth 1 2023-04-20T20:17:03 2023-04-20T20:17:03 2023-04-20T20:17:04 00:00:01 00:00:01 COMPLETED
venkata@t-green:/home/venkata/test/slurm_chain$ █

```

Output after execution

```

venkata@t-green:/home/venkata/test/slurm_chain$ cat first.txt
first job, job ID is 299
venkata@t-green:/home/venkata/test/slurm_chain$ cat j_s_300.out
Read Line: first job, job ID is 299
venkata@t-green:/home/venkata/test/slurm_chain$ █

```

# Individually submit an array job to cluster

## User: venkata

### Create new batch file

vi ind\_job\_arr.sh

### Add the following content

```
#!/bin/bash
#SBATCH --job-name=t_green_ind_venkata
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --time=00:01:00
#SBATCH --mem=1mb
#SBATCH --array=1-2
#SBATCH --output=j_%A_%a.out
#SBATCH --error=j_s_%j.err
cat /a/util/demo/slurm/run/slurm_arr/file_${SLURM_ARRAY_TASK_ID}.txt
```

### Execute the sbatch command

sbatch ind\_job\_arr.sh

```
venkata@t-green:/home/venkata$ vi ind_job_arr.sh
venkata@t-green:/home/venkata$ sbatch ind_job_arr.sh
Submitted batch job 310
venkata@t-green:/home/venkata$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0        up       infinite    4    idle i-hemanth,i-kush,i-roja,i-venkata
batch*     up       infinite    4    idle i-hemanth,i-kush,i-roja,i-venkata
```

```
venkata@t-green:/home/venkata$ cat j_310_
j_310_1.out  j_310_2.out
venkata@t-green:/home/venkata$ cat j_310_
j_310_1.out  j_310_2.out
```

Outputs are logged into the files: j\_%job%\_1.out, j\_%job%\_2.out

```
venkata@t-green:/home/venkata$ cat j_310_1.out
Here is file one

venkata@t-green:/home/venkata$ cat j_310_2.out
Here is file 2

venkata@t-green:/home/venkata$ █
```

# Individually submit an array job to cluster

## User: kush

### Create new batch file

```
vi ind_job_arr.sh
```

### Add the following content

```
#!/bin/bash
#SBATCH --job-name=t_green_ind_kush
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --time=00:01:00
#SBATCH --mem=1mb
#SBATCH --array=1-2
#SBATCH --output=j_%A_%a.out
#SBATCH --error=j_s_%j.err
cat /a/util/demo/slurm/run/slurm_arr/file_${SLURM_ARRAY_TASK_ID}.txt
```

### Execute the sbatch command

```
sbatch ind_job_arr.sh
```

**Outputs are logged into the files:** j\_%job%\_1.out, j\_%job%\_2.out

```
kush01@t-green:/home/kush$ vi ind_job_arr.sh
kush01@t-green:/home/kush$ ls
start_inst.sh  stop_inst.sh
kush01@t-green:/home/kush$ cd
kush01@t-green:/home/kush01$ vi ind_job_arr.sh
kush01@t-green:/home/kush01$ sbatch ind_job_arr.sh
Submitted batch job 312
kush01@t-green:/home/kush01$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0        up    infinite     4   idle i-hemanth,i-kush,i-roja,i-venkata
batch*     up    infinite     4   idle i-hemanth,i-kush,i-roja,i-venkata
kush01@t-green:/home/kush01$ cat j_312_
j_312_1.out  j_312_2.out
kush01@t-green:/home/kush01$ cat j_312_1.out
Here is file one
kush01@t-green:/home/kush01$ cat j_312_2.out
Here is file 2
kush01@t-green:/home/kush01$
```

# Individually submit an array job to cluster

**User: hemanth**

## Create new batch file

vi ind\_job\_arr.sh

## Add the following content

```
#!/bin/bash
#SBATCH --job-name=t_green_ind_hemanth
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --time=00:01:00
#SBATCH --mem=1mb
#SBATCH --array=1-2
#SBATCH --output=j_%A_%a.out
#SBATCH --error=j_s_%j.err
cat /a/util/demo/slurm/run/slurm_arr/file_${SLURM_ARRAY_TASK_ID}.txt
```

## Execute the sbatch command

sbatch ind\_job\_arr.sh

**Outputs are logged into the files:** j\_%job%\_1.out, j\_%job%\_2.out

```
[hemanth@t-green:/home/hemanth$ vi ind_job_arr.sh
[hemanth@t-green:/home/hemanth$ sbatch ind_job_arr.sh
Submitted batch job 314
[hemanth@t-green:/home/hemanth$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0       up    infinite     4    idle i-hemanth,i-kush,i-roja,i-venkata
batch*    up    infinite     4    idle i-hemanth,i-kush,i-roja,i-venkata
[hemanth@t-green:/home/hemanth$ cat j_314_
j_314_1.out j_314_2.out
[hemanth@t-green:/home/hemanth$ cat j_314_1.out
Here is file one
[hemanth@t-green:/home/hemanth$ cat j_314_2.out
Here is file 2
[hemanth@t-green:/home/hemanth$ █
```

# Individually submit an array job to cluster

**User: roja**

## Create new batch file

vi ind\_job\_arr.sh

## Add the following content

```
#!/bin/bash
#SBATCH --job-name=t_green_ind_roja
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --time=00:01:00
#SBATCH --mem=1mb
#SBATCH --array=1-2
#SBATCH --output=j_%A_%a.out
#SBATCH --error=j_s_%j.err
cat /a/util/demo/slurm/run/slurm_arr/file_${SLURM_ARRAY_TASK_ID}.txt
```

## Execute the sbatch command

sbatch ind\_job\_arr.sh

**Outputs are logged into the files:** j\_%job%\_1.out, j\_%job%\_2.out

```
roja@t-green:/home/roja$ vi ind_job_arr.sh
roja@t-green:/home/roja$ sbatch ind_job_arr.sh
Submitted batch job 316
roja@t-green:/home/roja$ sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ig0        up      infinite    4    idle i-hemanth,i-kush,i-roja,i-venkata
batch*     up      infinite    4    idle i-hemanth,i-kush,i-roja,i-venkata
roja@t-green:/home/roja$ cat j_316_
j_316_1.out j_316_2.out
roja@t-green:/home/roja$ cat j_316_1.out
Here is file one

roja@t-green:/home/roja$ cat j_316_2.out
Here is file 2

roja@t-green:/home/roja$ █
```

**Sub assignment 3:** All the team members have registered for Lucidcharts.

And have collaborated for the following diagram

[https://lucid.app/lucidchart/7f2494da-b82c-4474-87aa-53ef9354b0af/edit?viewport\\_loc=-102%2C-103%2C2581%2C1300%2CNSTQuOvS557b&invitationId=inv\\_2afb396e-faa1-4383-983f-855f4f08869b](https://lucid.app/lucidchart/7f2494da-b82c-4474-87aa-53ef9354b0af/edit?viewport_loc=-102%2C-103%2C2581%2C1300%2CNSTQuOvS557b&invitationId=inv_2afb396e-faa1-4383-983f-855f4f08869b)

